

Applications of metric evaluation

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON



Kevin Huo
Instructor

Four categories of outcomes

| | Actual Positives | Actual Negatives |
|-------------------------|----------------------------|----------------------------|
| Positive Predictions | True Positives (TP) | False Positives (FP) |
| Negative Predictions | False Negatives (FN) | True Negatives (TN) |

- First part of category (true/false) represents whether model was correct or not
- Second part of the category (positive/negative) represents the target label the model applied

Interpretations of four categories

- If model predicts there is a click, then there is a bid for that impression which costs money
- If no click predicted, no bidding and hence no cost
- True positives (TP): money gained (impressions paid for that were clicked on).
- False positives (FP): money lost (impressions that were paid for, but not clicked).
- True negatives (TN): money saved (no click predicted so no impressions bought).
- False negatives (FN): money lost out on (no click predicted, but would have been actual click in reality).

Confusion matrix

```
print(confusion_matrix(y_test, y_pred))
```

```
[[8163  166]
 [1517  154]]
```

```
# Order: tn, fp, fn, tp
print(confusion_matrix(y_test, y_pred).ravel())
```

```
[8163, 166, 1517, 154]
```

ROI analysis

- Assume: some cost c and return r per X number of impressions

```
total_return = tp * r
```

```
total_cost = (tp + fp) * c
```

```
tp * r > (tp + fp) * c
```

```
roi = total_return / total_spent
```

Let's practice!

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON

Model evaluation

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON



Kevin Huo
Instructor

Precision and recall

- Precision: proportion of clicks relative to total number of impressions, $TP / (TP + FP)$
 - Higher precision means higher ROI on ad spend
- Recall: the proportion of clicks gotten of all clicks available, $TP / (TP + FN)$
 - Higher recall means better targeting of relevant audience

Calculating precision and recall

```
print(precision_score(  
    y_test, y_pred, average = 'weighted'))
```

```
0.73
```

```
print(recall_score(  
    y_test, y_pred, average = 'weighted'))
```

```
0.75
```

Baseline classifiers

- It is important to evaluate classifiers relative to an appropriate baseline
 - The baseline here, due to imbalanced nature of click data, is a classifier that always predicts no click

```
y_pred = np.asarray([0 for x in range(len(X_test))])
```

```
[[0]  
 [0] ...]
```

Implications on ROI analysis

- For the baseline classifier, `tp` and `fp` will be zero
- Therefore total return and total spend will be zero, and ROI undefined
- Confusion matrix via `confusion_matrix()` along with `ravel()` to get the four categories of outcomes

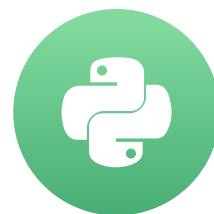
```
total_return = tp * r
total_spent = (tp + fp) * cost
roi = total_return / total_spent
```

Let's practice!

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON


Tuning models

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON



Kevin Huo
Instructor

Regularization

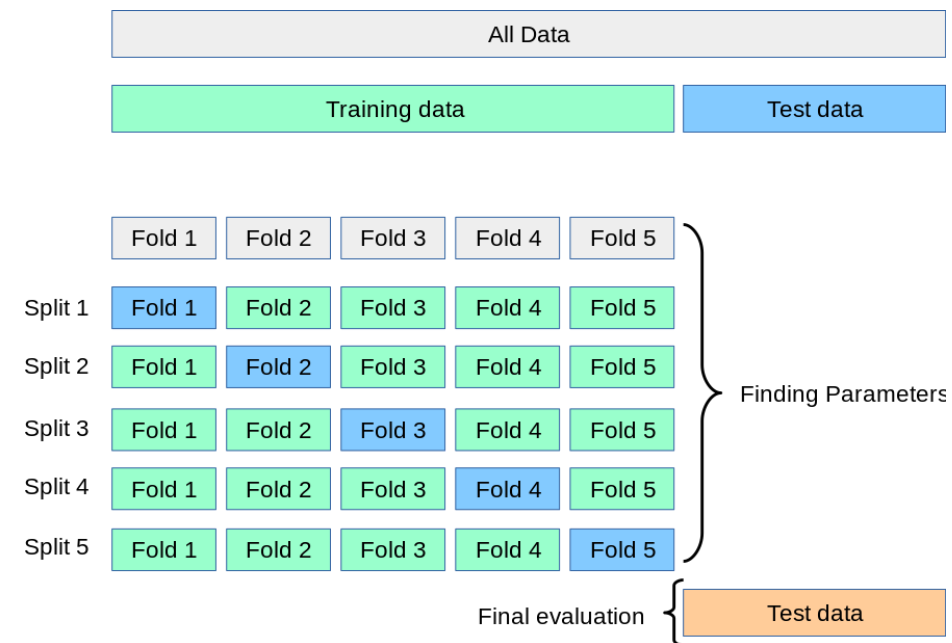


- Regularization: addressing overfitting by altering the magnitude of coefficients of parameters within a model
- Regularization can increase performance metrics and hence ROI on ad spend

Examples of regularization

- Logistic Regression: the `C` parameter is the inverse of the regularization strength.
- From least to most complex: `C=0.05 < C=0.5 < C=1`
- Decision Tree: the `max_depth` parameter controls how many layers deep the tree can grow.
- From least to most complex: `max_depth=3 < max_depth=5 < max_depth=10`

Cross validation



- For each of the k folds, that fold will be used as a testing set (for validation) while other $k-1$ are used as training.
- Therefore, you have k evaluations of model performance.
- Note you still have the separate evaluation testing set.

Examples of cross validation

```
k_fold = KFold(n_splits = 4, random_state = 0)
```

```
for i in [3, 5, 10]:  
    clf = DecisionTreeClassifier(max_depth = i)  
    cv_precision = cross_val_score(  
        clf, X_train, y_train, cv = k_fold,  
        scoring = 'precision_weighted')
```

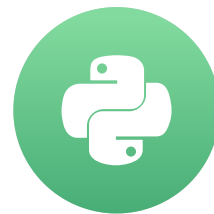
- Scoring strings: precision_weighted, recall_weighted, roc_auc

Let's practice!

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON

Ensembles and hyperparameter tuning


PREDICTING CTR WITH MACHINE LEARNING IN PYTHON



Kevin Huo
Instructor

Ensemble methods

“Bagging” : **B**ootstrap **AGG**regat**ING**



- Bagging: random samples selected for different models, then models are individually trained and combined.

Random forests

```
clf = RandomForestClassifier()  
print(clf)
```

```
RandomForestClassifier(  
    bootstrap=True,  
    ...  
    max_depth = 10,  
    ...  
    n_estimators = 100,  
    ...)
```

Hyperparameter tuning

- Hyperparameter: parameters configured before training, and external to a model
- Examples of parameters but NOT hyperparameters: slope coefficient in linear regression, weights in logistic regression, etc.
- Examples of hyperparameters: `max_depth` , `n_estimators` , etc.

Grid search

```
param_grid = {'n_estimators': n_estimators,  
              'max_depth': max_depth}
```

```
clf = GridSearchCV(estimator = model,  
                   param_grid = param_grid,  
                   scoring = 'roc_auc')  
  
print(clf.best_score_)  
print(clf.best_estimator_)
```

```
0.6777
```

```
RandomForestClassifier(max_depth = 100, ...)
```

Let's practice!

PREDICTING CTR WITH MACHINE LEARNING IN PYTHON